

Rindrit Telaku

Data Engineer

✉ rindritelaku@gmail.com 📍 Pristina 🔗 [linkedin.com/in/rindrittelaku](https://www.linkedin.com/in/rindrittelaku) 🐙 github.com/RindTel
🌐 rindrittelaku.vercel.app

PROFESSIONAL EXPERIENCE

Junior Ingestion Specialist, ProQu

Remote | 05/2026 – Present

- **Built a config driven AWS ingestion framework** using PySpark, EMR Serverless, and Apache Iceberg, allowing new external data sources to be onboarded primarily through configuration.
- **Developed a Bronze/Silver/Gold lakehouse architecture** with Apache Iceberg, incremental merges, and schema evolution, transforming heterogeneous external data into analytics-ready datasets.
- **Orchestrated production pipelines** with AWS Step Functions and EventBridge Scheduler, with infrastructure managed through Terraform and failure monitoring via SNS and CloudWatch.
- **Implemented declarative data quality gates** covering freshness, null rates, uniqueness, and custom SQL checks, alongside **800+ pytest tests** for ingestion and transformation pipelines.
- **Contributed to the FastAPI/PostgreSQL backend**, delivering hybrid search, an AWS Bedrock LLM assistant, news alerting, and role based access control.

PROJECTS

Backtesting Engine 🌐, Python, Polars, MLflow, Streamlit

Present

- **Built a lookahead safe backtesting engine** with history isolated strategy APIs and next bar execution, preventing strategies from accessing future market data.
- **Implemented automated leakage testing** that corrupts post cutoff market data and verifies that pre cutoff trading decisions remain unchanged.
- **Modelled realistic execution costs** including fees, bid ask spread, square root market impact, and volume based fill limits, enabling strategy capacity analysis across position sizes.
- **Added reproducible experiment tracking** with MLflow, dataset content hashes, in/out of sample splits, and automated parameter sweeps, backed by **116 tests and GitHub Actions CI**.

RAG System 🌐, LangChain, FAISS, Ollama, Qwen 2.5

- **Built a fully offline document Q&A system** with local LLM inference, persistent document indexing, source citations, and no external API dependencies.
- **Implemented hybrid retrieval** combining BM25 and vector search with Reciprocal Rank Fusion and cross encoder reranking, improving evaluation pass rate from **52% to 97%** on a custom 31 question evaluation set.
- **Engineered retrieval evaluation and ablation testing**, measuring retrieval hit rate, context recall, and context window overflows across multiple configurations.
- **Packaged the system as a cross platform desktop application** using llama.cpp and ONNX based embeddings, with GitHub Actions validating packaged builds through automated end to end self tests.

Transit Lens 🌐, DuckDB, dbt, Streamlit

- **Built a multi city GTFS analytics pipeline** processing **31M+ transit stop events** across three European networks into a local DuckDB warehouse.
- **Implemented incremental, idempotent ingestion** using SHA256 file fingerprinting, with automatic processing of new feeds and protection against duplicate loads.
- **Developed dbt staging and mart layers** to clean and model transit data into route activity, busiest stops, and hourly service patterns, handling GTFS specific data quirks and ID collisions.
- **Built an interactive Streamlit dashboard** with map based stop visualization and automated refresh triggered by new feed ingestion.

SKILLS

Data Engineering: Python, SQL, PySpark, Polars, dbt, DuckDB, PostgreSQL

Cloud & Infrastructure: AWS (S3, EMR Serverless, Athena, Step Functions, EventBridge, CloudWatch), Terraform, Git, GitHub Actions

AI/ML: LLMs, RAG, FAISS, BM25, Vector Search, MLflow

Visualization: Streamlit, Altair, PyDeck

EDUCATION

BSc in Computer Science and Engineering,
UBT - University for Business and Technology 🌐

2024 – Present